

## **Model Training and Evaluation**

### **1. Compiling the Model**

Compiling a model involves configuring its learning process. During this step, you specify:

- **Optimizer:** Algorithm used to adjust the weights, e.g., adam.
- **Loss Function:** Function minimized during training, e.g., `sparse_categorical_crossentropy` for classification.
- **Metrics:** Metrics monitored during training and evaluation, e.g., accuracy.

This configuration prepares the model for training by setting up the required tools to update weights and measure performance.

### **2. Fitting the Model**

Fitting the model means training it on the provided data. Key parameters include:

- **Training Data (`x_train`, `y_train`):** Inputs and corresponding labels.
- **Epochs:** Number of times the learning algorithm will work through the entire training dataset.
- **Batch Size:** Number of samples processed before the model's internal parameters are updated.
- **Validation Data (`x_val`, `y_val`):** Data used to evaluate the model's performance during training.

The fit method trains the model by iteratively adjusting weights to minimize the loss function.

### **3. Evaluating the Model**

Evaluating the model involves assessing its performance on a separate test dataset. Key outputs are:

- **Loss:** Final value of the loss function on the test data.
- **Accuracy (or other metrics):** Performance metric(s) on the test data.

The evaluate method provides an unbiased estimate of the model's performance by measuring how well it generalizes to new data.

### **4. Making Predictions**

Making predictions involves using the trained model to infer outputs for new input data. This is done using the predict method, which outputs the model's predictions for the provided inputs.

### **Summary**

1. **Compile:** Configure the model with an optimizer, loss function, and metrics.
2. **Fit:** Train the model on the training data, optionally validating on separate validation data.
3. **Evaluate:** Measure the model's performance on a test dataset.
4. **Predict:** Use the trained model to generate predictions for new data.

These steps form the core workflow for training and evaluating machine learning models using Keras, providing a structured approach to developing, validating, and deploying predictive models.

Configuring the model for training using the `compile()` method.

Configuring a model for training using the `compile()` method in Keras is a crucial step in the machine learning workflow. This method sets up the model with the necessary components for training: the optimizer, the loss function, and the evaluation metrics. Below, we'll delve into each component and how to configure them.

```
import tensorflow as tf
```

```
from tensorflow.keras import layers, models
```

```
# Define the model
```

```
model = models.Sequential([
    layers.Conv2D(32, (3, 3), activation='relu', input_shape=(64, 64, 3)), # Convolutional layer
    layers.MaxPooling2D((2, 2)), # Max pooling layer
    layers.Conv2D(64, (3, 3), activation='relu'), # Another convolutional layer
    layers.MaxPooling2D((2, 2)), # Another max pooling layer
    layers.Flatten(), # Flatten the output to 1D
    layers.Dense(64, activation='relu'), # Fully connected layer
    layers.Dense(10, activation='softmax') # Output layer with 10 classes
])
```

```
# Compile the model
```

```
model.compile(optimizer='adam', # Optimizer
              loss='sparse_categorical_crossentropy', # Loss function
              metrics=['accuracy']) # Evaluation metric
```

```
# Model summary
```

```
model.summary()
```

Model Definition:

- **Conv2D:** Applies 2D convolution with 32 filters of size 3x3, using ReLU activation. It processes input images of shape (64, 64, 3).
- **MaxPooling2D:** Reduces the spatial dimensions by taking the maximum value in each 2x2 patch.
- **Flatten:** Converts the 2D feature maps into a 1D vector.

- Dense: Fully connected layers, where the final layer has 10 units with softmax activation for classification into 10 classes.

Compiling the Model:

- optimizer='adam': Uses the Adam optimization algorithm, which adjusts learning rates based on the gradients.
- loss='sparse\_categorical\_crossentropy': Computes the cross-entropy loss for multi-class classification problems where labels are integers.
- metrics=['accuracy']: Tracks accuracy during training and evaluation to measure performance.

```

/usr/local/lib/python3.10/dist-packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not pass an `input_shape` to
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
Model: "sequential_1"

```

Layer (type)	Output Shape	Param #
conv2d_2 (Conv2D)	(None, 62, 62, 32)	896
max_pooling2d_2 (MaxPooling2D)	(None, 31, 31, 32)	0
conv2d_3 (Conv2D)	(None, 29, 29, 64)	18,496
max_pooling2d_3 (MaxPooling2D)	(None, 14, 14, 64)	0
flatten_6 (Flatten)	(None, 12544)	0
dense_4 (Dense)	(None, 64)	802,880
dense_5 (Dense)	(None, 10)	650

Total params: 822,922 (3.14 MB)  
 Trainable params: 822,922 (3.14 MB)  
 Non-trainable params: 0 (0.00 B)

Training models using the fit() method.

The fit() method in Keras is used to train a model on your data. It takes care of running the training loop, adjusting model weights based on the optimizer and loss function, and optionally monitoring performance on validation data. Here's a breakdown of how to use it, along with an example.

Parameters of fit()

1. x: Input data.
2. y: Target data (labels).
3. batch\_size: Number of samples per gradient update.
4. epochs: Number of times the entire dataset is passed forward and backward through the neural network.
5. validation\_data: Data on which to evaluate the loss and any model metrics at the end of each epoch.

6. verbose: Verbosity mode (0, 1, or 2).

```
import tensorflow as tf
from tensorflow.keras import layers, models
import numpy as np

# Define the model
model = models.Sequential([
    layers.Conv2D(32, (3, 3), activation='relu', input_shape=(64, 64, 3)),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Flatten(),
    layers.Dense(64, activation='relu'),
    layers.Dense(10, activation='softmax')
])

# Compile the model
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

# Generate example data
x_train = np.random.random((1000, 64, 64, 3))
y_train = np.random.randint(10, size=(1000,))
x_val = np.random.random((200, 64, 64, 3))
y_val = np.random.randint(10, size=(200,))

# Train the model
history = model.fit(x_train, y_train, # Training data
                   batch_size=32, # Number of samples per gradient update
```

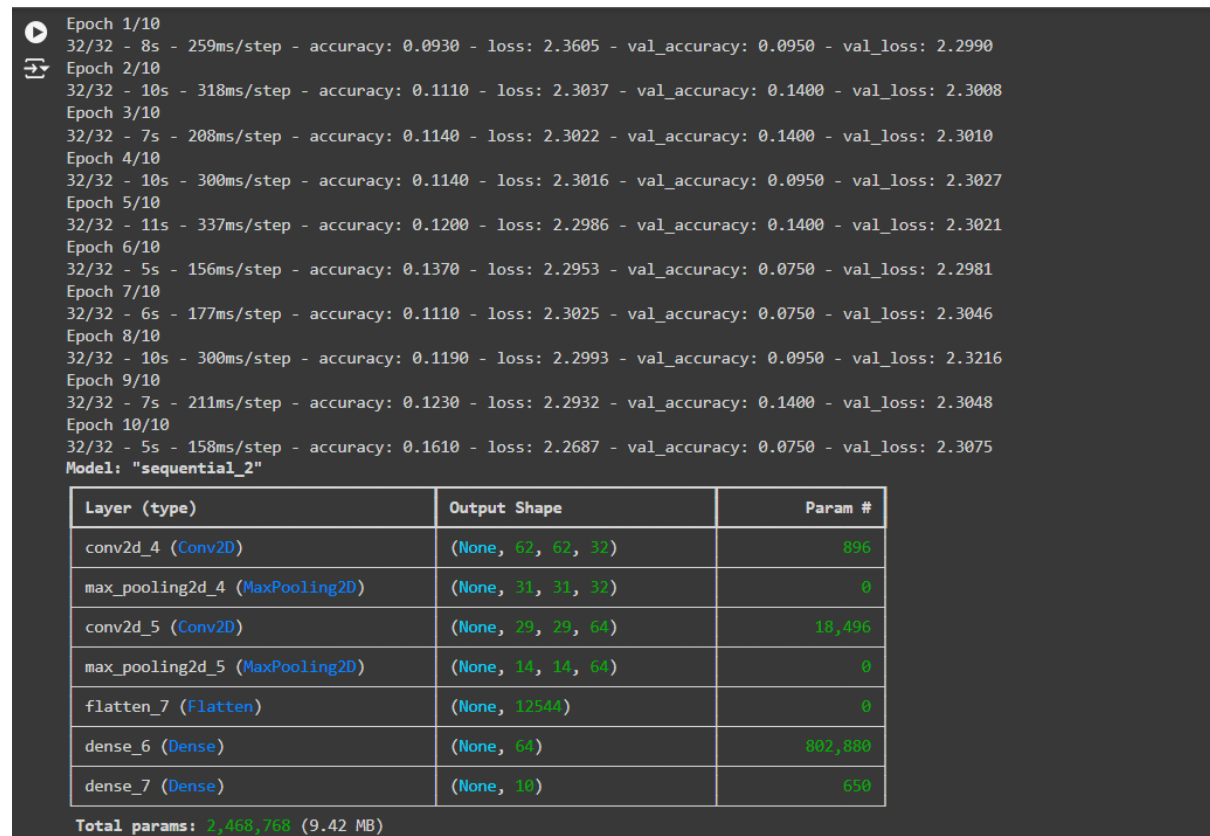
```
epochs=10, # Number of epochs to train

validation_data=(x_val, y_val), # Validation data

verbose=2) # Verbosity mode
```

# Model summary

```
model.summary()
```



Epoch 1/10  
32/32 - 8s - 259ms/step - accuracy: 0.0930 - loss: 2.3605 - val\_accuracy: 0.0950 - val\_loss: 2.2990  
Epoch 2/10  
32/32 - 10s - 318ms/step - accuracy: 0.1110 - loss: 2.3037 - val\_accuracy: 0.1400 - val\_loss: 2.3008  
Epoch 3/10  
32/32 - 7s - 208ms/step - accuracy: 0.1140 - loss: 2.3022 - val\_accuracy: 0.1400 - val\_loss: 2.3010  
Epoch 4/10  
32/32 - 10s - 300ms/step - accuracy: 0.1140 - loss: 2.3016 - val\_accuracy: 0.0950 - val\_loss: 2.3027  
Epoch 5/10  
32/32 - 11s - 337ms/step - accuracy: 0.1200 - loss: 2.2986 - val\_accuracy: 0.1400 - val\_loss: 2.3021  
Epoch 6/10  
32/32 - 5s - 156ms/step - accuracy: 0.1370 - loss: 2.2953 - val\_accuracy: 0.0750 - val\_loss: 2.2981  
Epoch 7/10  
32/32 - 6s - 177ms/step - accuracy: 0.1110 - loss: 2.3025 - val\_accuracy: 0.0750 - val\_loss: 2.3046  
Epoch 8/10  
32/32 - 10s - 300ms/step - accuracy: 0.1190 - loss: 2.2993 - val\_accuracy: 0.0950 - val\_loss: 2.3216  
Epoch 9/10  
32/32 - 7s - 211ms/step - accuracy: 0.1230 - loss: 2.2932 - val\_accuracy: 0.1400 - val\_loss: 2.3048  
Epoch 10/10  
32/32 - 5s - 158ms/step - accuracy: 0.1610 - loss: 2.2687 - val\_accuracy: 0.0750 - val\_loss: 2.3075  
Model: "sequential\_2"

Layer (type)	Output Shape	Param #
conv2d_4 (Conv2D)	(None, 62, 62, 32)	896
max_pooling2d_4 (MaxPooling2D)	(None, 31, 31, 32)	0
conv2d_5 (Conv2D)	(None, 29, 29, 64)	18,496
max_pooling2d_5 (MaxPooling2D)	(None, 14, 14, 64)	0
flatten_7 (Flatten)	(None, 12544)	0
dense_6 (Dense)	(None, 64)	802,880
dense_7 (Dense)	(None, 10)	650

Total params: 2,468,768 (9.42 MB)  
Trainable params: 803,426 (3.14 MB)

Setting batch size, epochs, and validation split.

In Keras, you can control several aspects of the training process using the `fit()` method, including batch size, epochs, and validation split. Here's a detailed explanation of how to set each of these parameters:

#### Batch Size

- **Definition:** The number of samples processed before the model's internal parameters are updated. A smaller batch size means the model is updated more frequently, but it may be less stable. A larger batch size processes more samples in one go, but updates less frequently.
- **Setting:** You can specify the batch size in the `fit()` method.

## Epochs

- Definition: The number of times the entire training dataset is passed forward and backward through the model. Each epoch involves a complete pass through the training data, with weights being updated after each batch.
- Setting: You specify the number of epochs in the fit() method.

## Validation Split

- Definition: The fraction of the training data to be used as validation data. This is a way to monitor the model's performance on data it hasn't seen during training, helping to detect overfitting.
- Setting: You provide the fraction of data to be used for validation directly in the fit() method.

```
import tensorflow as tf
```

```
from tensorflow.keras import layers, models
```

```
import numpy as np
```

```
# Define the model
```

```
model = models.Sequential([  
    layers.Conv2D(32, (3, 3), activation='relu', input_shape=(64, 64, 3)),  
    layers.MaxPooling2D((2, 2)),  
    layers.Conv2D(64, (3, 3), activation='relu'),  
    layers.MaxPooling2D((2, 2)),  
    layers.Flatten(),  
    layers.Dense(64, activation='relu'),  
    layers.Dense(10, activation='softmax')  
])
```

```
# Compile the model
```

```
model.compile(optimizer='adam',  
              loss='sparse_categorical_crossentropy',  
              metrics=['accuracy'])
```

```
# Generate example data
```

```
x_train = np.random.random((1000, 64, 64, 3))
```

```

y_train = np.random.randint(10, size=(1000,))

# Train the model with batch size, epochs, and validation split

history = model.fit(x_train, y_train, # Training data

                    batch_size=32, # Number of samples per gradient update

                    epochs=10, # Number of epochs to train

                    validation_split=0.2, # Fraction of training data to use as validation data

                    verbose=2) # Verbosity mode

# Model summary

model.summary()

```

```

/usr/local/lib/python3.10/dist-packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not pass an `input_shape`/`input_dim` argument
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
Epoch 1/5
25/25 - 6s - 236ms/step - accuracy: 0.1050 - loss: 2.3829 - val_accuracy: 0.1050 - val_loss: 2.3068
Epoch 2/5
25/25 - 6s - 236ms/step - accuracy: 0.1100 - loss: 2.3014 - val_accuracy: 0.1050 - val_loss: 2.3032
Epoch 3/5
25/25 - 10s - 388ms/step - accuracy: 0.1375 - loss: 2.2977 - val_accuracy: 0.1050 - val_loss: 2.3062
Epoch 4/5
25/25 - 5s - 197ms/step - accuracy: 0.1063 - loss: 2.2976 - val_accuracy: 0.1100 - val_loss: 2.2999
Epoch 5/5
25/25 - 4s - 179ms/step - accuracy: 0.1238 - loss: 2.2973 - val_accuracy: 0.1100 - val_loss: 2.3042
Model: "sequential"

```

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 62, 62, 32)	896
max_pooling2d (MaxPooling2D)	(None, 31, 31, 32)	0
conv2d_1 (Conv2D)	(None, 29, 29, 64)	18,496
max_pooling2d_1 (MaxPooling2D)	(None, 14, 14, 64)	0
flatten (Flatten)	(None, 12544)	0
dense (Dense)	(None, 64)	802,800
dense_1 (Dense)	(None, 10)	650

```

Total params: 2,468,768 (9.42 MB)
Trainable params: 822,922 (3.14 MB)
Non-trainable params: 0 (0.00 B)
Optimizer params: 1,645,846 (6.28 MB)

```

Evaluating model performance using the evaluate() method.

The evaluate() method in Keras is used to assess the performance of a trained model on a given dataset. This method returns the loss value and any additional metrics specified during model compilation. It provides insight into how well the model performs on unseen data or on a validation set.

Parameters of evaluate()

1. x: Input data for evaluation.
2. y: True labels corresponding to the input data.
3. batch\_size: Number of samples per gradient update. If not specified, it defaults to 32.
4. verbose: Verbosity mode (0, 1, or 2). Controls the amount of information printed during evaluation.

```

import tensorflow as tf

from tensorflow.keras import layers, models

import numpy as np

# Define the model
model = models.Sequential([
    layers.Conv2D(32, (3, 3), activation='relu', input_shape=(64, 64, 3)),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Flatten(),
    layers.Dense(64, activation='relu'),
    layers.Dense(10, activation='softmax')
])

# Compile the model
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

# Generate example training data
x_train = np.random.random((1000, 64, 64, 3))
y_train = np.random.randint(10, size=(1000,))

# Train the model
model.fit(x_train, y_train, batch_size=32, epochs=10, validation_split=0.2, verbose=2)

# Generate example test data
x_test = np.random.random((200, 64, 64, 3))
y_test = np.random.randint(10, size=(200,))

```



```
# Evaluate the model on test data

test_loss, test_accuracy = model.evaluate(x_test, y_test, verbose=2)

print(f'Test Loss: {test_loss}')
print(f'Test Accuracy: {test_accuracy}')
```