

# Pandas

## Data Cleaning

Data cleaning means fixing bad data in your data set.

Bad data could be:

- Empty cells
- Data in wrong format
- Wrong data
- Duplicates
- Empty Cells
- Empty cells can potentially give you a wrong result when you analyze data.
- Remove Rows
- One way to deal with empty cells is to remove rows that contain empty cells.
- This is usually OK, since data sets can be very big, and removing a few rows will not have a big impact on the result.

Python code

data.csv

```
import pandas as pd

df = pd.read_csv('data.csv')

new_df = df.dropna()

print(new_df.to_string())

#Notice in the result that some rows have been removed (row 18, 22 and 28).

#These rows had cells with empty values.
```

	Duration	Date	Pulse	Maxpulse	Calories
0	60	'2020/12/01'	110	130	409.1
1	60	'2020/12/02'	117	145	479.0
2	60	'2020/12/03'	103	135	340.0
3	45	'2020/12/04'	109	175	282.4
4	45	'2020/12/05'	117	148	406.0
5	60	'2020/12/06'	102	127	300.0
6	60	'2020/12/07'	110	136	374.0
7	450	'2020/12/08'	104	134	253.3
8	30	'2020/12/09'	109	133	195.1
9	60	'2020/12/10'	98	124	269.0
10	60	'2020/12/11'	103	147	329.3
11	60	'2020/12/12'	100	120	250.7
12	60	'2020/12/12'	100	120	250.7
13	60	'2020/12/13'	106	128	345.3

**Note:** By default, the `dropna()` method returns a *new* DataFrame, and will not change the original.

If you want to change the original DataFrame, use the `inplace = True` argument:

Example

Remove all rows with NULL values:

```
import pandas as pd
```

```
df = pd.read_csv('data.csv')
```

```
df.dropna(inplace = True)
```

```
print(df.to_string())
```

**Note:** Now, the `dropna(inplace = True)` will NOT return a new DataFrame, but it will remove all rows containing NULL values from the original DataFrame.

# Pandas

## Replace Empty Values

Another way of dealing with empty cells is to insert a *new* value instead.

This way you do not have to delete entire rows just because of some empty cells.

The fillna() method allows us to replace empty cells with a value:

Example

Replace NULL values with the number 130:

```
import pandas as pd
```

```
df = pd.read_csv('data.csv')
```

```
df.fillna(130, inplace = True)
```

## Replace Only For Specified Columns

Python code

data.csv

```
import pandas as pd

df = pd.read_csv('data.csv')

df["Calories"].fillna(130, inplace = True)

print(df.to_string())

#This operation Inserts 130 in empty cells in the "Calories" column (row 18 and 28).
```

	Duration	Date	Pulse	Maxpulse	Calories
0	60	'2020/12/01'	110	130	409.1
1	60	'2020/12/02'	117	145	479.0
2	60	'2020/12/03'	103	135	340.0
3	45	'2020/12/04'	109	175	282.4
4	45	'2020/12/05'	117	148	406.0
5	60	'2020/12/06'	102	127	300.0
6	60	'2020/12/07'	110	136	374.0
7	450	'2020/12/08'	104	134	253.3
8	30	'2020/12/09'	109	133	195.1
9	60	'2020/12/10'	98	124	269.0
10	60	'2020/12/11'	103	147	329.3
11	60	'2020/12/12'	100	120	250.7
12	60	'2020/12/12'	100	120	250.7
13	60	'2020/12/13'	106	128	345.3
14	60	'2020/12/14'	104	132	379.3
15	60	'2020/12/15'	98	123	275.0
16	60	'2020/12/16'	98	120	215.2

## Replace Using Mean, Median, or Mode

A common way to replace empty cells, is to calculate the mean, median or mode value of the column.

Pandas uses the mean() median() and mode() methods to calculate the respective values for a specified column:

Example

**Calculate the MEAN, and replace any empty values with it:**

```
import pandas as pd
```

```
df = pd.read_csv('data.csv')
```

# Pandas

```
x = df["Calories"].mean()
```

```
df["Calories"].fillna(x, inplace = True)
```

Python code	data.csv																																																																																																
<pre>import pandas as pd  df = pd.read_csv('data.csv')  x = df["Calories"].mean()  df["Calories"].fillna(x, inplace = True)  print(df.to_string())  #As you can see in row 18 and 28, the empty values from "Calories" was replaced with the mean: 304.68</pre>	<table><tr><th></th><th>Duration</th><th>Date</th><th>Pulse</th><th>Maxpulse</th><th>Calories</th></tr><tr><td>0</td><td>60</td><td>'2020/12/01'</td><td>110</td><td>130</td><td>409.10</td></tr><tr><td>1</td><td>60</td><td>'2020/12/02'</td><td>117</td><td>145</td><td>479.00</td></tr><tr><td>2</td><td>60</td><td>'2020/12/03'</td><td>103</td><td>135</td><td>340.00</td></tr><tr><td>3</td><td>45</td><td>'2020/12/04'</td><td>109</td><td>175</td><td>282.40</td></tr><tr><td>4</td><td>45</td><td>'2020/12/05'</td><td>117</td><td>148</td><td>406.00</td></tr><tr><td>5</td><td>60</td><td>'2020/12/06'</td><td>102</td><td>127</td><td>300.00</td></tr><tr><td>6</td><td>60</td><td>'2020/12/07'</td><td>110</td><td>136</td><td>374.00</td></tr><tr><td>7</td><td>450</td><td>'2020/12/08'</td><td>104</td><td>134</td><td>253.30</td></tr><tr><td>8</td><td>30</td><td>'2020/12/09'</td><td>109</td><td>133</td><td>195.10</td></tr><tr><td>9</td><td>60</td><td>'2020/12/10'</td><td>98</td><td>124</td><td>269.00</td></tr><tr><td>10</td><td>60</td><td>'2020/12/11'</td><td>103</td><td>147</td><td>329.30</td></tr><tr><td>11</td><td>60</td><td>'2020/12/12'</td><td>100</td><td>120</td><td>250.70</td></tr><tr><td>12</td><td>60</td><td>'2020/12/12'</td><td>100</td><td>120</td><td>250.70</td></tr><tr><td>13</td><td>60</td><td>'2020/12/13'</td><td>106</td><td>128</td><td>345.30</td></tr><tr><td>14</td><td>60</td><td>'2020/12/14'</td><td>104</td><td>132</td><td>379.30</td></tr></table>		Duration	Date	Pulse	Maxpulse	Calories	0	60	'2020/12/01'	110	130	409.10	1	60	'2020/12/02'	117	145	479.00	2	60	'2020/12/03'	103	135	340.00	3	45	'2020/12/04'	109	175	282.40	4	45	'2020/12/05'	117	148	406.00	5	60	'2020/12/06'	102	127	300.00	6	60	'2020/12/07'	110	136	374.00	7	450	'2020/12/08'	104	134	253.30	8	30	'2020/12/09'	109	133	195.10	9	60	'2020/12/10'	98	124	269.00	10	60	'2020/12/11'	103	147	329.30	11	60	'2020/12/12'	100	120	250.70	12	60	'2020/12/12'	100	120	250.70	13	60	'2020/12/13'	106	128	345.30	14	60	'2020/12/14'	104	132	379.30
	Duration	Date	Pulse	Maxpulse	Calories																																																																																												
0	60	'2020/12/01'	110	130	409.10																																																																																												
1	60	'2020/12/02'	117	145	479.00																																																																																												
2	60	'2020/12/03'	103	135	340.00																																																																																												
3	45	'2020/12/04'	109	175	282.40																																																																																												
4	45	'2020/12/05'	117	148	406.00																																																																																												
5	60	'2020/12/06'	102	127	300.00																																																																																												
6	60	'2020/12/07'	110	136	374.00																																																																																												
7	450	'2020/12/08'	104	134	253.30																																																																																												
8	30	'2020/12/09'	109	133	195.10																																																																																												
9	60	'2020/12/10'	98	124	269.00																																																																																												
10	60	'2020/12/11'	103	147	329.30																																																																																												
11	60	'2020/12/12'	100	120	250.70																																																																																												
12	60	'2020/12/12'	100	120	250.70																																																																																												
13	60	'2020/12/13'	106	128	345.30																																																																																												
14	60	'2020/12/14'	104	132	379.30																																																																																												

## Data of Wrong Format

Cells with data of wrong format can make it difficult, or even impossible, to analyze data.

To fix it, you have two options: remove the rows, or convert all cells in the columns into the same format.

Pandas has a `to_datetime()` method for this:

Python code

data.csv

```
import pandas as pd

df = pd.read_csv('data.csv')

df['Date'] = pd.to_datetime(df['Date'])

print(df.to_string())
```

	Duration	Date	Pulse	Maxpulse	Calories
0	60	2020-12-01	110	130	409.1
1	60	2020-12-02	117	145	479.0
2	60	2020-12-03	103	135	340.0
3	45	2020-12-04	109	175	282.4
4	45	2020-12-05	117	148	406.0
5	60	2020-12-06	102	127	300.0
6	60	2020-12-07	110	136	374.0
7	450	2020-12-08	104	134	253.3
8	30	2020-12-09	109	133	195.1
9	60	2020-12-10	98	124	269.0
10	60	2020-12-11	103	147	329.3
11	60	2020-12-12	100	120	250.7
12	60	2020-12-12	100	120	250.7

## Removing Rows

The result from the converting in the example above gave us a NaT value, which can be handled as a NULL value, and we can remove the row by using the `dropna()` method.

Example

Remove rows with a NULL value in the "Date" column:

```
df.dropna(subset=['Date'], inplace = True)
```

## Wrong Data

"Wrong data" does not have to be "empty cells" or "wrong format", it can just be wrong, like if someone registered "199" instead of "1.99".

Replace Values:

```
df.loc[7, 'Duration'] = 45
```

# Pandas

Python code

data.csv

```
import pandas as pd

df = pd.read_csv('data.csv')

for x in df.index:
    if df.loc[x, "Duration"] > 120:
        df.loc[x, "Duration"] = 120

print(df.to_string())
```

	Duration	Date	Pulse	Maxpulse	Calories
0	60	'2020/12/01'	110	130	409.1
1	60	'2020/12/02'	117	145	479.0
2	60	'2020/12/03'	103	135	340.0
3	45	'2020/12/04'	109	175	282.4
4	45	'2020/12/05'	117	148	406.0
5	60	'2020/12/06'	102	127	300.0
6	60	'2020/12/07'	110	136	374.0
7	120	'2020/12/08'	104	134	253.3
8	30	'2020/12/09'	109	133	195.1
9	60	'2020/12/10'	98	124	269.0
10	60	'2020/12/11'	103	147	329.3
11	60	'2020/12/12'	100	120	250.7
12	60	'2020/12/12'	100	120	250.7
13	60	'2020/12/13'	106	128	345.3

## Removing Rows

Python code

data.csv

```
import pandas as pd

df = pd.read_csv('data.csv')

for x in df.index:
    if df.loc[x, "Duration"] > 120:
        df.drop(x, inplace = True)

#remember to include the 'inplace = True' argument to make the changes in the original DataFrame object
#instead of returning a copy

print(df.to_string())
```

	Duration	Date	Pulse	Maxpulse	Calories
0	60	'2020/12/01'	110	130	409.1
1	60	'2020/12/02'	117	145	479.0
2	60	'2020/12/03'	103	135	340.0
3	45	'2020/12/04'	109	175	282.4
4	45	'2020/12/05'	117	148	406.0
5	60	'2020/12/06'	102	127	300.0
6	60	'2020/12/07'	110	136	374.0
8	30	'2020/12/09'	109	133	195.1
9	60	'2020/12/10'	98	124	269.0
10	60	'2020/12/11'	103	147	329.3
11	60	'2020/12/12'	100	120	250.7
12	60	'2020/12/12'	100	120	250.7
13	60	'2020/12/13'	106	128	345.3
14	60	'2020/12/14'	104	132	379.3
15	60	'2020/12/15'	98	123	275.0
16	60	'2020/12/16'	98	120	215.2
17	60	'2020/12/17'	100	120	200.0

## Removing Duplicates:

To discover duplicates, we can use the duplicated() method.

The duplicated() method returns a Boolean values for each row:

Example [Get your own Python Server](#)

Returns True for every row that is a duplicate, otherwise False:

```
print(df.duplicated())
```

Python code

data.csv

```
import pandas as pd

df = pd.read_csv('data.csv')

df.drop_duplicates(inplace = True)

print(df.to_string())

#Notice that row 12 has been removed from the result
```

	Duration	Date	Pulse	Maxpulse	Calories
0	60	'2020/12/01'	110	130	409.1
1	60	'2020/12/02'	117	145	479.0
2	60	'2020/12/03'	103	135	340.0
3	45	'2020/12/04'	109	175	282.4
4	45	'2020/12/05'	117	148	406.0
5	60	'2020/12/06'	102	127	300.0
6	60	'2020/12/07'	110	136	374.0
7	450	'2020/12/08'	104	134	253.3
8	30	'2020/12/09'	109	133	195.1
9	60	'2020/12/10'	98	124	269.0
10	60	'2020/12/11'	103	147	329.3
11	60	'2020/12/12'	100	120	250.7
13	60	'2020/12/13'	106	128	345.3

## Data Correlations:

Finding Relationships

# Pandas

A great aspect of the Pandas module is the `corr()` method.

The `corr()` method calculates the relationship between each column in your data set

Python code

data.csv

```
import pandas as pd

df = pd.read_csv('data.csv')

print(df.corr())
```

	Duration	Pulse	Maxpulse	Calories
Duration	1.000000	-0.059452	-0.250033	0.344341
Pulse	-0.059452	1.000000	0.269672	0.481791
Maxpulse	-0.250033	0.269672	1.000000	0.335392
Calories	0.344341	0.481791	0.335392	1.000000

## Plotting

Pandas uses the `plot()` method to create diagrams.

We can use Pyplot, a submodule of the Matplotlib library to visualize the diagram on the screen.


Import pyplot from Matplotlib and visualize our DataFrame:

```
import pandas as pd
import matplotlib.pyplot as plt
```

```
df = pd.read_csv('data.csv')
```

```
df.plot()
```

```
plt.show()
```



## Scatter Plot

Specify that you want a scatter plot with the `kind` argument:

```
kind = 'scatter'
```

A scatter plot needs an x- and a y-axis.

# Pandas

In the example below we will use "Duration" for the x-axis and "Calories" for the y-axis.

Include the x and y arguments like this:

```
x = 'Duration', y = 'Calories'
```


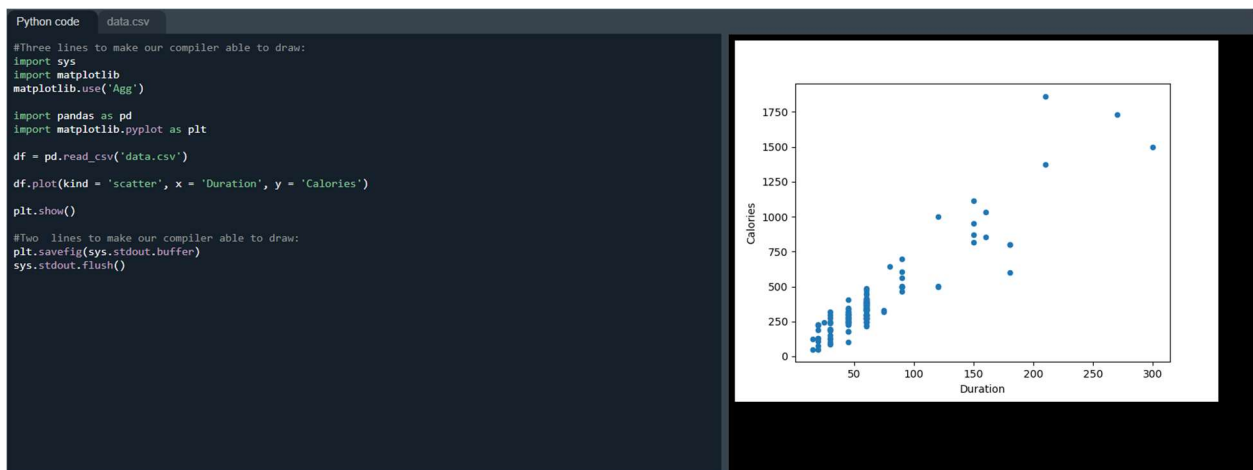
Example

```
import pandas as pd
import matplotlib.pyplot as plt
```

```
df = pd.read_csv('data.csv')
```

```
df.plot(kind = 'scatter', x = 'Duration', y = 'Calories')
```

```
plt.show()
```



## Histogram

# Pandas

Use the kind argument to specify that you want a histogram:

```
kind = 'hist'
```

A histogram needs only one column.

A histogram shows us the frequency of each interval, e.g. how many workouts lasted between 50 and 60 minutes?

In the example below we will use the "Duration" column to create the histogram:

Example

```
df["Duration"].plot(kind = 'hist')
```

```
#Three lines to make our compiler able to draw:
import sys
import matplotlib
matplotlib.use('Agg')


import pandas as pd
import matplotlib.pyplot as plt

df = pd.read_csv('data.csv')

df["Duration"].plot(kind = 'hist')

plt.show()

#Two lines to make our compiler able to draw:
plt.savefig(sys.stdout.buffer)
sys.stdout.flush()
```



Thank you